

Building tidy R packages for the OMOP common data model

Edward Burn, Martí Català

2025-10-08

Table of contents

Preface	1
I. Getting started	3
1. Scope	5
1.1. Is this package needed?	5
2. Ecosystem	6
2.1. How the package fits in the ecosystem	6
3. Contributing to a package	7
3.1. GitHub issues	7
3.1.1. Opening GitHub issues	7
3.1.2. What makes a good issue	7
3.2. Responding to GitHub issues	8
3.3. Contributing code	10
Contribute documentation	10
Fixing a bug	10
Adding new functionality	10
3.4. Opening pull requests	11
3.4.1. Initial checks	12
3.4.2. Run tests	12
3.4.3. Adhere to code style	13
3.4.4. Run check() before opening a pull request	13
3.4.5. Opening the pull request	13

Table of contents

3.5. Reviewing pull requests	14
4. Set up your environment	15
4.1. Create an empty package	16
4.2. Add a license	16
4.3. Create a first mock function	18
4.4. Check the package	18
4.5. Create README.Rmd	18
4.6. Create website	18
4.7. Set up GitHub	18
4.8. Package template	18
 II. Principles	 19
5. Core dependencies	20
5.1. omopgenerics	20
5.1.1. Classes	20
5.1.2. Methods	20
5.1.3. Input validation functions	20
5.1.4. Manipulate a <code>cdm_reference</code> object	20
5.1.5. Manipulate a <code>summarised_result</code> object	20
5.2. Other useful packages	20
6. Function interfaces	22
7. Conventions	23
7.1. Casing	23
7.2. Argument naming	23
8. Function outputs	24
9. Error messages and input validation	25
9.1. omopgenerics validation functions	25
9.2. omopgenerics assert functions	28

Table of contents

9.3. Examples	28
9.4. Conclusions	30
10. Styling your code	31
11. Reexporting from other packages	32
 III. Contributing documentation	 33
12. Documenting code	34
13. Documenting functions	35
13.1. Title	35
13.2. Description	36
13.3. Parameters	36
13.4. Returns	37
13.5. Export	37
13.6. Examples	38
14. Readme	39
15. Vignettes	40
16. Documentation website	41
 IV. Adding unit tests	 42
17. Unit testing	43
18. Testing across multiple systems	44
19. Efficient testing	45
20. Documenting tests	46

Table of contents

V. GitHub	47
21. Set up GitHub	48
22. Automatic R-CDM checks	49
23. Website deployment	50
24. Test coverage	51
25. Testing multiple DBMS	52
VI. Maintaing a package	53
26. Submit to CRAN	54

Preface

This is an **opinionated** book on how we develop packages working in the Tidy R OMOP CDM space.

Our book, although covers some parts of basic R package development assumes the user knows the very basis. If you are new to R package development we would recommend you the following book: R Packages; Learn how to create a package, the fundamental unit of shareable, reusable, and reproducible R code.

The book is divided in four parts:

- **(I)** Getting started: covers how to start with an empty package, and helps you think about which is the scope of your package or guide you on how to contribute to existing packages.
- **(II)** Principles: this is the most important part of the book as it covers which packages to use, the different naming conventions, classes to use, how to validate the inputs and more.
- **(III)** Documentation: covers how to create the documentation of the packages, such as populate readme, documenting functions, documenting code, vignettes and more.
- **(IV)** Testing: covers how to write tests for the packages, what to test and how to do it efficiently.
- **(V)** GitHub: covers how to set up automatic GitHub testing and documentation using GitHub actions.

Preface

- **(VI)** Maintaining a package: once a package is ready for a first release you will have to maintain it, this packages covers how to submit to cran for first time and maintain the package for the future.

Part I.

Getting started

In this first part it is covered how to ...

1. Scope

1.1. Is this package needed?

Before deciding to develop a package you must ask yourself many questions starting with: *Is this package needed?* A package must cover need and be refined.

2. Ecosystem

2.1. How the package fits in the ecosystem

3. Contributing to a package

3.1. GitHub issues

3.1.1. Opening GitHub issues

The first step to contributing to an existing package is by opening issues on its GitHub repository. These issues could be about bugs you encountered when using the package, requests for additional functionality (that you might even want to add yourself), or clarification questions on package documentation.

3.1.2. What makes a good issue

If reporting a bug, then a reprex makes the maintainers life much, much easier (and in turn increases the likelihood of a quick fix being introduced much higher!). This reprex will allow the developer to quickly reproduce the issue. This is already halfway to solving it, given the old adage > “A problem well stated is half solved.”

Often an issue may arise when running code against a database with real patient-level data which can make it challenging to quickly create a reprex (as the patient data for which the bug was seen cannot be shared). In this situation it may be possible to reproduce the problem on an available synthetic dataset like Eunomia, in which case the package developer will be able to reproduce the bug with the synthetic data. However, it may be that you only encounter the issue for some particular set of patients

3. *Contributing to a package*

with specific characteristics which are not seen in the synthetic data. In this case you can, for example, use the `omock` package to create a set of synthetic patients with these characteristics so that the developer can use this same data to reproduce the problem.

Although of course it would be ideal to include a `reprex`, this may not always be possible. In such cases please do make sure to provide a bug report, with as much information as possible, as this will still be extremely useful and appreciated.

If opening an issue to ask for clarifications on documentation or with requests for additional functionality, the more precise and specific the issue the better. If you have a long laundry list of feature requests, it would generally be better to open each as separate issues so they can be addressed incrementally.

Note, even if you are not the maintainer of package but want to open a pull request to change documentation or add functionality then try to always open an issue to discuss this first. Even some seemingly trivial changes might be out of scope for the package (or they might have already been added in development branches of the repo). However well-intentioned, “drive-by pull requests” are often challenging for package maintainers to deal with.

3.2. **Responding to GitHub issues**

On the other side of the equation if you are the maintainer of package, receiving issues is one the main ways you will receive feedback on the package. Get used to starting with somewhat vague issues that will need some back and forth to get to the route of the problem or request for additional functionality.

When a bug is fixed in the package as a rule you should also add a corresponding test to make sure that it will stay fixed into the future. One

3. Contributing to a package

benefit of being given a reprex is that not only will it help you to fix the problem but the reprex itself can be used as the test for the package. By receiving and acting on issues raised, which will often be somewhat exotic edge cases or people using the package in some way that you hadn't expected, the package will become hardened over time.

When it comes to requests for adding functionality, or even offers to add it, this is a tricky balancing act. Remember that you will be the one maintaining the package in the long-term. It is often the case that the quick addition of something seemingly useful can make a package more difficult to maintain or extend in the future. Typically it is worth waiting a day or two at least before agreeing to add something to a package so that you don't rush into making an addition or change that you will later come to regret.

Another point to keep in mind when considering issues with requests for functionality is whether to add more general functionality than is being asked for. A well-intentioned user may be asking for something very specific to be added, but often underlying this request will be a broader need for more general functionality. This may tie in to other requests received, or may even need more discussion with users to better understand whether there is a broader need at play. general solutions

Remember that sadly few people be opening issues with praise if they've used your package successfully. The majority of users will use the code you've written without a problem and likely without even realising how much work has gone into it. And the people that do open issues are providing an extremely valuable service that will lead to the software getting better and better over time. So try to keep this in mind when tackling the seemingly never-ending list of issues you'll get if your software is being used.

3.3. Contributing code

Contribute documentation

Contributing documentation is one of the best ways to improve a package. Fixing typos, adding clarifications, and even writing whole vignettes will almost always be highly welcomed by the maintainer and improve the experience of other users. You can find issues related to documentation searching for the label *documentation*. Also you can open an issue if you think that something is not documented properly and propose the maintainer to fix it yourself.

You can read more details on how the documentation of the packages is done in the relevant chapters.

Fixing a bug

Contributing code is a more involved task. Before starting, make sure you've at the very least interacted with the maintainer so that they will be in favour of the changes or additions you are planning. You can do that replying to the issue you want to fix. Apart from fixing the bug, make sure to add tests alongside code changes to show that it is truly fixed. When fixing a bug existing tests should typically be unchanged (unless there is an issue with the test itself). If you are having to change existing tests to fix a bug this is normally a sign that there might be some more profound issue and will likely need back and forth between you and the maintainer for this to be addressed.

Adding new functionality

If adding functionality this again should be accompanied with tests to show that it adds the desired behaviour. Moreover, updates to package docu-

3. Contributing to a package

mentation, including changes or additions to vignettes might be needed. Expect back and forth with the package maintainer, as they might well have feedback on the implementation you propose.

3.4. Opening pull requests

Changes to packages should come via pull requests. Create a branch or a fork of the code, make your change, and open your pull request. If opening a pull request, more than anything try to create a pull request that will be easy for the package maintainer to review and merge. A pull request will be better for everyone involved when it addresses just one specific issue and affects less than 250 lines of code. One quick way to create complexity for the maintainer is to open a pull request that addresses multiple issues at the same time. Although this could be more efficient for the person opening the pull request, it often slows down the review by increasing the complexity of the review for the maintainer.

Before starting to contribute any code, first make sure the package tests are all passing. To do so after cloning run the following code:

```
devtools::check()
```

The output should be:

```
0 errors v | 0 warnings v | 0 notes v
```

If not raise an issue before going any further (although please first make sure you have all the packages from imports and suggests installed).

Now you are ready to do your code contribution. Add the relevant code and when you are happy with the changes that you have made, please follow these steps to open a pull request:

3. Contributing to a package

3.4.1. Initial checks

Run the below to update and check package documentation:

```
devtools::document()  
devtools::check_man()
```

Test that the examples work:

```
devtools::run_examples()
```

Test that the readme and vignettes work and update the results if the output of any function has changed:

```
devtools::build_readme()  
devtools::build_vignettes()
```

3.4.2. Run tests

As you then contribute code, make sure that all the current tests and any you add continue to pass. All package tests can be run together with:

```
devtools::test()
```

Code to add new functionality should be accompanied by tests. Code coverage can be checked using:

```
# note, you may first have to detach the package  
# detach("package:IncidencePrevalence", unload=TRUE)  
devtools::test_coverage()
```

3. Contributing to a package

3.4.3. Adhere to code style

Please adhere to the code style when adding any new code. Do not though restyle any code unrelated to your pull request as this will make code review more difficult.

```
lintr::lint_package(linters = lintr::linters_with_defaults(  
  lintr::object_name_linter(styles = "camelCase")  
))
```

3.4.4. Run `check()` before opening a pull request

Before opening any pull request please make sure to run:

```
devtools::check()
```

Please make sure that the output is the expected:

```
0 errors v | 0 warnings v | 0 notes v
```

Any error or warning will make your pull request actions fail. Although notes can pass the github action tests we encourage to fix all the notes as this is a requirement to keep the package on cran.

3.4.5. Opening the pull request

Once you've made sure that checks are passing and you are happy with your code additions you can open the pull request. When opening the pull request you must take into account the following:

3. Contributing to a package

- Write a meaningful title, titles of pull requests are later used to document the changes done. Make sure that the title of the pull request describes the issue that is fixed, therefore the documentation of the changes will be easier.
- Link the issue that your pull request is closing (we strongly encourage to fix issues one by one and do not include multiple issues in the same pull request). You can link a pull request and an issue using any of the gitub closing words or you can do it manually in the *development* section (bottom of the right sidebar).
- Describe any potential issue you want to remark. Help the maintainer review your pull request indicating if there is anything you are not sure about or that you want some feedback on. Don't hesitate to ask any question if needed.

3.5. Reviewing pull requests

Reviewing pull requests is a very important step that the maintainers have to do. When reviewing a pull request you have to find a good equilibrium between being kind (we want to encourage people to contribute to our packages) and being strict (you -as a maintainer- will be the ultimate responsible for that code; so review it carefully and only accept if you are happy about it and you feel comfortable maintaining that code).

4. Set up your environment

To build an R package you will need R installed and an IDE: preferably R Studio or Positron.

You will need to install devtools and usethis as they are key to follow the different steps to get started.

```
# use base R
install.packages(c("usethis", "devtools"))

# or pak (recommended)
library(pak)
pak_install(c("usethis", "devtools"))
```

To work with the Tidy R OMOP ecosystem, you will also need to install the omopgenerics package. Alongside with other packages that may be useful:

- PatientProfiles for data manipulation.
- CodelistDiagnostics to query the vocabularies.
- visOmopResults for data visualisation.

You can see the list of Tidy R packages in our website.

4. Set up your environment

4.1. Create an empty package

To create an empty package you can do it with *usethis*:

```
library(usethis)

create_package(path = "path/to/folder")
```

💡 Package name

The package will be named by the name of the folder.

4.2. Add a license

Adding licenses is very important so you are protected and users know how they can use your product. In general we release using the Apache 2.0 license.

You can add a license easily with *usethis*:

```
use_apl2_license()
```

This will add the `LICENSE.md` file and populate the `DESCRIPTION` file with the relevant license field.

ℹ License agreements

License agreements supported by *usethis* are:

- More permissive:
 - MIT: simple and permissive.

4. Set up your environment

- Apache 2.0: MIT + provides patent protection.
- Copyleft:
 - GPL v2: requires sharing of improvements.
 - GPL v3: requires sharing of improvements.
 - AGPL v3: requires sharing of improvements.
 - LGPL v2.1: requires sharing of improvements.
 - LGPL v3: requires sharing of improvements.
- Creative commons licenses appropriate for data packages:
 - CC0: dedicated to public domain.
 - CC-BY: Free to share and adapt, must give appropriate credit.

```
use_mit_license(copyright_holder = NULL)

use_gpl_license(version = 3, include_future = TRUE)

use_agpl_license(version = 3, include_future = TRUE)

use_lgpl_license(version = 3, include_future = TRUE)

use_apache_license(version = 2, include_future = TRUE)

use_cc0_license()

use_ccby_license()

use_proprietary_license(copyright_holder = "...")
```

4. Set up your environment

4.3. Create a first mock function

4.4. Check the package

4.5. Create README.Rmd

4.6. Create website

4.7. Set up GitHub

GitHub actions are very useful to ensure the continuous integration workflow. The fifth part of the book is dedicated to GitHub and GitHub Actions at that stage you will learn how to use github for:

- Automatic checking your package
- Automatic deployment of a documentation website for your package
- Automatic assessment of test coverage
- Automatic testing across multiple sources

4.8. Package template

Part II.

Principles

5. Core dependencies

The *Tidy R OMOP CDM* packages rely on `dplyr` and the `tidyverse` packages to manipulate the data from the `cdm` object. The `cdm` object is an object that contains all the tables available and is central for the data manipulation. This `cdm` object is defined in `omopgenerics`.

5.1. `omopgenerics`

5.1.1. Classes

5.1.2. Methods

5.1.3. Input validation functions

5.1.4. Manipulate a `cdm_reference` object

5.1.5. Manipulate a `summarised_result` object

5.2. Other useful packages

The *omopverse* has some core packages that will be useful to simplify the scope of your package:

5. Core dependencies

- PatientProfiles: this package is very useful for data manipulation, you can find extract basic demographic information, intersections of cohorts and summarise data in a standard format. Take a look at the list of functions in PatientProfiles to prevent you from duplicating code.
- visOmopResults: this package is very useful to create data visualisations (tables and plots) from standardised data. There are many different supported types of tables (`r visOmopResults::tableType()`) and with little customisation you can get quite pretty tables. For plots the support is a bit more limited but it creates some simple ggplot2 visualisations with very little effort. Take a look at the list of functions in visOmopResults to prevent you from duplicating code.
- CodelistGenerator: this package can be useful to query the vocabularies. Take a look at the list of functions in CodelistGenerator to prevent you from duplicating code.
- omock: this package can be very useful for testing. With very few lines of code you can create a mock CDM object with your desired specifications, specially useful can be the function `mockCdmFromTables()` that will create a viable CDM object from the supplied information. Take a look at the list of functions in omock to prevent you from duplicating code. This package is quite useful, but usually listed in Suggests as you will only need this package for testing purposes.

6. Function interfaces

```
generate...CohortSet, ...Cohort, add..., summarise..., plot...,  
table...,
```

7. Conventions

Our code adheres to some style conventions as described in this chapter.

7.1. Casing

In general we use **camelCase** naming for arguments and functions.

7.2. Argument naming

8. Function outputs

9. Error messages and input validation

Giving users good and informative error messages is key for a good user experience. To do so it is important that we perform an input validation at the beginning of each function. On the other hand we do not want to spend more time checking the input than executing the function, so don't overdo it. In general we would check that any input function has the desired type and length and that the desired evaluation looks feasible with the current input parameters.

9.1. omopgenerics validation functions

Some arguments that are consistent across different functions and packages have their own *validate* functions in the **omopgenerics** package:

- `validateAgeGroupArgument()` is used to validate the `ageGroup` argument. The output `ageGroup` will always be formatted as a named list (name of the age group) and each age group will be defined by named intervals.
- `validateCdmArgument()` is used to validate the `cdm` argument. By default only the class is validated as this can take time, but specific checks can be triggered if needed.
- `validateCohortArgument()` is used to validate `cohort` argument, very used in many packages, this validates that the input is a properly formatted cohort. Extra checks can be triggered if needed.

9. Error messages and input validation

- `validateCohortIdArgument()` is used to validate `cohortId` argument
- `validateConceptSetArgument()`
- `validateNameArgument()`
- `validateNameStyle()`
- `validateResultArgument()`
- `validateStrataArgument()`
- `validateWindowArgument()`

It is important that we assign the output to the variable as the object might change during the validation process to ensure different allowed inputs but as the output of the validation process will always be the same this simplifies the code as you do not have to think about the different allowed inputs.

the `ageGroup` argument can be a good example of this behavior:

```
library(omopgenerics, warn.conflicts = FALSE)
ageGroup <- validateAgeGroupArgument(ageGroup = list(c(0, 1), c(10, 20)))
ageGroup
```

```
$age_group
$age_group$`0 to 1`
[1] 0 1
```

```
$age_group$`10 to 20`
[1] 10 20
```

```
ageGroup <- validateAgeGroupArgument(ageGroup = list(
  my_column = list("young" = c(0, 19), 20, c(21, Inf)),
  list(c(0, 9), c(10, 19), c(20, 29), c(30, Inf))
))
ageGroup
```

9. Error messages and input validation

```
$my_column  
$my_column$young  
[1] 0 19
```

```
$my_column$`20 to 20`  
[1] 20 20
```

```
$my_column$`21 or above`  
[1] 21 Inf
```

```
$age_group_2  
$age_group_2$`0 to 9`  
[1] 0 9
```

```
$age_group_2$`10 to 19`  
[1] 10 19
```

```
$age_group_2$`20 to 29`  
[1] 20 29
```

```
$age_group_2$`30 or above`  
[1] 30 Inf
```

As you can see the output is always a named list that contains named intervals the function itself will also throw explanatory errors if they are not properly formatted:

```
validateAgeGroupArgument(  
  ageGroup = list(age_group1 = list(c(0, 19), c(20, Inf)), age_group2 = list(c(0, In  
  multipleAgeGroup = FALSE  
)
```

Error:

9. Error messages and input validation

! Multiple age group are not allowed

```
validateAgeGroupArgument(  
  ageGroup = list(age_group1 = list(c(-5, 19), c(20, Inf)))  
)
```

```
Error in `purrr::map()`:  
i In index: 1.  
i With name: age_group1.  
Caused by error:  
! Elements of `ageGroup` argument must be greater or equal to "0".
```

```
validateAgeGroupArgument(  
  ageGroup = NULL, null = FALSE  
)
```

```
Error:  
! `ageGroup` argument can not be NULL.
```

9.2. omopgenerics assert functions

The omopgenerics package contains some functions for simple validation steps this can be useful helpers to validate an input with a single line of code, they also contain arguments to check if they have

9.3. Examples

Let's say we have a function with four arguments (`cohort`, `cohortId`, `window` and `overlap`), we could easily validate the input arguments of the function with 4 lines of code:

9. Error messages and input validation

```
myFunction <- function(cohort, cohortId = NULL, window = c(0, Inf), overlap = FALSE)
# input check
cohort <- omopgenerics::validateCohortArgument(cohort = cohort)
cohortId <- omopgenerics::validateCohortIdArgument(cohortId = {{cohortId}}, cohort
window <- omopgenerics::validateWindowArgument(window = window)
omopgenerics::assertLogical(overlap, length = 1)

# code ...

}
```

Note the `{{` symbols are needed to be able to use tidyselect verbs such as `starts_with()` or `contains()`.

A second example that needs some custom extra code can be:

```
myFunction <- function(cdm, conceptSet, days = 180L, startDate = NULL, overlap = TRUE)
# input check
cdm <- omopgenerics::validateCdmArgument(cdm = cdm)
conceptSet <- omopgenerics::validateCdmArgument(conceptSet = conceptSet)
omopgenerics::assertNumeric(days, integerish = TRUE, min = 0, length = 1)
omopgenerics::assertDate(startDate, length = 1, null = TRUE)
if (overlap & days > 365) {
  cli::cli_abort(c(x = "{.var days} is can not be >= 365 if {.var overlap} is TRUE
}

# code ...

}
```

You can throw custom error and warning messages using the **cli** package.

9.4. Conclusions

Validating arguments is a very important step to give user a good experience and prevent running undesired code in big datasets. The **omop-generics** provides you with some functionality to keep the validation step short and consistent with other packages.

10. Styling your code

11. Reexporting from other packages

Part III.

Contributing documentation

12. Documenting code

13. Documenting functions

We use the `roxygen2` package for documenting our functions. Using this package will help us to write the `.Rd` files in the `man` directory which is what will be rendered when users call `?my_function` or `help("my_function")` and peruse a package website.

Above each exported function we'll create a roxygen **block**, with lines starting with `#'`. The general structure is shown below and we'll go through each in turn.

```
#' Title
#'\n
#' @description
#'\n
#' @param x
#' @param y
#'\n
#' @returns
#' @export
#'\n
#' @examples
```

13.1. Title

We want to provide a concise yet informative title for our function. The title should typically be written in sentence case but without any full stop

13. Documenting functions

at the end.

If we run `help(package = "mypackage")` we can see all the functions and their titles for a package. From this we can quickly see how informative our titles are.

13.2. Description

Every function should have a description. This could be a sentence or two for a simple function, or multiple paragraphs for a more complex function. The description should summarise the purpose of the function and explain any internal assumptions or decisions made to implement it.

When writing the description we can use markdown commands, for example to add bullet points, italics, etc. This will often help make longer descriptions more readable.

13.3. Parameters

After our description of the function as a whole, we will then describe each of the inputs to the function. In general we should be able to describe an argument in a few sentences. We should use them to tell the user what is allowed as an input and then what the input will be used for.

If we're working on a package where we have multiple arguments with the same argument we can document arguments just once. We can create a helper like below for parameter `x`.

```
#' Helper for consistent documentation of `x`.  
#'  
#' @param x  
#'
```

13. Documenting functions

```
#' @name xDoc  
#' @keywords internal
```

We can then reuse this documentation across the package by using `@inheritParams`

```
#' Title  
#'  
#' @description  
#'  
#' @inheritParams xDoc  
#' @param y  
#'  
#' @returns  
#' @export  
#'  
#' @examples
```

13.4. Returns

We must also document what the function will return. For example, we can tell a user that our function will return a tibble with a certain set of columns. If our function has side-effects then we should also document these.

13.5. Export

By including the `@export` tag our documentation will be generated as an `.Rd` file and users will see it in the package website, etc. If instead we want to document an internal function then we can replace this with `@noRd`.

13.6. Examples

Last but by no means least are a set of examples where we provide R code that shows how the function can be used in practice. These examples are quite possibly the most important piece of documentation, as it is many users will look at these before reading any of the above.

Typically we should try to give a number of informative examples that show users how the function can be used and provides some intuition on the impact of different arguments.

14. Readme

- badges

15. Vignettes

16. Documentation website

Part IV.

Adding unit tests

17. Unit testing

do not rely on mock data

18. Testing across multiple systems

19. Efficient testing

20. Documenting tests

Part V.

GitHub

21. Set up GitHub

22. Automatic R-CDM checks

23. Website deployment

24. Test coverage

25. Testing multiple DBMS

Part VI.

Maintaing a package

26. Submit to CRAN

- What version to release (major, minor or patch)
- Forced releases by cran
- news
- reverse dependency